

Lecture 24: PG Vector Store Intro

Why Not a Simple Vector Store?

- Simple Vector Store is good for learning and experiments but not ideal for production.
- Real projects need a persistent and reliable vector database.
- That's where PGVector comes in.

What is PGVector?

- PGVector is an open-source PostgreSQL extension for handling vector embeddings.
- It allows storing embeddings in Postgres and performing similarity searches.
- Built directly on PostgreSQL, making it reliable and easy to integrate if you already use Postgres

Installation & Setup

- PostgreSQL does not include PGVector by default.
- You need to install/enable it separately.
- Options to set it up:
 - Docker Image/Docker Compose is the easiest approach.
 - Local installation (platform-specific steps for Windows, Linux, and Mac).

Spring AI + PGVector:

- In Spring AI, PGVector works as a Vector Store implementation.
- You will need:
 - Dependencies → add the pgvector dependency along with Spring AI.
 - Database Configuration → URL, username, password.
 - Vector Store Properties → specify:
 - index type
 - distance type (cosine, L2, etc.)
 - dimensions
 - max document batch size

Key Takeaways:

- PGVector lets you store and query embeddings inside PostgreSQL.
- Best to use in production where persistence and scalability matter.